

Practice Exam 3 Solution

STAT 3113 Regression Analysis

The Boston Housing Dataset

In R, the MASS library contains Boston data set, which has 506 rows and 14 columns, which records medv (median house value) and 13 predictors for 506 neighborhoods around Boston. The data frame contains the following columns:

- crim —per capita crime rate by town.
- zn — proportion of residential land zoned for lots over 25,000 sq.ft.
- indus — proportion of non-retail business acres per town.
- chas — Charles River dummy variable (= 1 if tract bounds river; 0 otherwise).
- nox — nitrogen oxides concentration (parts per 10 million).
- rm — average number of rooms per dwelling.
- age — proportion of owner-occupied units built prior to 1940.
- dis — weighted mean of distances to five Boston employment centres.
- rad — index of accessibility to radial highways.
- tax — full-value property-tax rate per \$10,000.
- ptratio — pupil-teacher ratio by town.
- black — $1000(Bk - 0.63)^2$ where Bk is the proportion of blacks by town.
- lstat — lower status of the population (percent).
- medv — median value of owner-occupied homes in \$1000s. (**dependent variable**)

Load the Data

```
library(MASS)
names(Boston)
```

```
## [1] "crim"    "zn"      "indus"   "chas"    "nox"     "rm"      "age"
## [8] "dis"     "rad"     "tax"     "ptratio" "black"   "lstat"   "medv"
```

Question 1

Conduct a stepwise regression analysis of the Boston data using R to find the “best” predictors of medv. Please use 0.15 for both α -to-remove and α -to-enter. Comment on the output.

```
library(olsrr)
fit_complete<-lm(medv~., data=Boston)
ols_step_both_p(fit_complete, details=FALSE, pent=0.15, prem = 0.15)
```

```
##
##                               Stepwise Selection Summary
```

##	Step	Variable	Added/ Removed	R-Square	Adj. R-Square	C(p)	AIC	RMSE
##	1	lstat	addition	0.544	0.543	362.7530	3288.9750	6.2158
##	2	rm	addition	0.639	0.637	185.6470	3173.5423	5.5403
##	3	ptratio	addition	0.679	0.677	111.6490	3116.0973	5.2294
##	4	dis	addition	0.690	0.688	91.4850	3099.3590	5.1386
##	5	nox	addition	0.708	0.705	59.7540	3071.4386	4.9939
##	6	chas	addition	0.716	0.712	47.1750	3059.9390	4.9326
##	7	black	addition	0.722	0.718	37.0590	3050.4384	4.8818
##	8	zn	addition	0.727	0.722	30.6240	3044.2750	4.8474
##	9	crim	addition	0.729	0.724	28.4180	3042.1546	4.8326
##	10	rad	addition	0.734	0.729	20.2660	3034.0687	4.7895
##	11	tax	addition	0.741	0.735	10.1150	3023.7264	4.7362

Answer:

- Response variable – medv
- Possible predictors – 13 the other variables.

From the R output, the best predictors of medv are: lstat, rm, ptratio, dis, nox, chas, black, zn, crim, rad, tax.

Question 2

What are the dangers associated with drawing inferences from the stepwise model?

Answer: The model was selected as the best model after many different models were fit to the data. There are two main dangers associate with using and interpreting this model:

1. The overall probability with making at least one Type I error is extremely high.
2. This model only includes linear components. Before using this model, both quadratic and interaction relationships in the model should be considered.

Question 3

Use all-possible-regressions-selection to find the “best” predictors of medv. Include the adjusted r-square plot and Cp plot, and illustrate how the choice is made.

```
library(leaps)
all.possible<-regsubsets(medv~., data=Boston, nvmax=13)

bestsubset.summary<-summary(all.possible)
bestsubset.summary$outmat
```

```
##          crim zn  indus chas nox rm  age dis rad tax ptratio black lstat
## 1  ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 2  ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 3  ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 4  ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 5  ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 6  ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 7  ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 8  ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 9  ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 10 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 11 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 12 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 13 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
```

```
rsq = bestsubset.summary$rsq
adjrsq = bestsubset.summary$adjr2
rss = bestsubset.summary$rss
cp = bestsubset.summary$cp
```

```
p=seq(1, 13, 1)
cbind(p, adjrsq)
```

```
##          p      adjrsq
## [1,] 1 0.5432418
## [2,] 2 0.6371245
## [3,] 3 0.6767036
## [4,] 4 0.6878351
## [5,] 5 0.7051702
## [6,] 6 0.7123567
## [7,] 7 0.7182560
## [8,] 8 0.7222072
## [9,] 9 0.7252743
## [10,] 10 0.7299149
## [11,] 11 0.7348058
## [12,] 12 0.7343282
## [13,] 13 0.7337897
```

```
cbind(p, cp)
```



```
##          p      cp
## [1,] 1 362.75295
## [2,] 2 185.64743
## [3,] 3 111.64889
## [4,] 4  91.48526
## [5,] 5  59.75364
## [6,] 6  47.17537
## [7,] 7  37.05889
## [8,] 8  30.62398
## [9,] 9  25.86592
## [10,] 10 18.20493
## [11,] 11 10.11455
## [12,] 12 12.00275
## [13,] 13 14.00000
```

```
cbind(p, rss)
```


```
##      p      rss
## [1,]  1 19472.38
## [2,]  2 15439.31
## [3,]  3 13727.99
## [4,]  4 13228.91
## [5,]  5 12469.34
## [6,]  6 12141.07
## [7,]  7 11868.24
## [8,]  8 11678.30
## [9,]  9 11526.12
## [10,] 10 11308.58
## [11,] 11 11081.36
## [12,] 12 11078.85
## [13,] 13 11078.78
```

```
### Plot the rsq, adjrsq, rss, cp vs p
```

```
par(mfrow=c(1,2))
plot(p, rsq)
plot(p, adjrsq)
```



```
plot(p, cp)
abline(1,1)
plot(p, rss)
```



```
par(mfrow=c(1,1))
```

Answer: The larger the value of R_a^2 , the better. On the other hand, we prefer models with a small value of C_p and a value of C_p near $p + 1$. Based on these criterion, we prefer the model with 11 predictors.

The 11 predictors are: lstat, rm, ptratio, dis, nox, chas, black, zn, crim, rad, tax.

Question 4

Compare the results in Question 1 and 3, which independent variables consistently are selected as the “best” predictors?

Answer: Based on the Question 1 and 3, the 11 independent variables are consistently chosen as the best predictors are: lstat, rm, ptratio, dis, nox, chas, black, zn, crim, rad, tax.


Question 5

Fit the first-order linear model with 10 independent variables: crim, chas, nox, rm, dis, rad, tax, ptratio, black, lstat. Plot the residual plots. Comment on the four residual plots one by one on the issues you found.

What model adjustments would you recommend?

```
fit = lm(medv~crim + chas + nox + rm + dis + rad + tax + ptratio + black + lstat, data = Boston)

par(mfrow=c(2,2))
plot(fit)
```




```
par(mfrow=c(1,1))
```

Answer: From the residual plot in the upper left corner, we see a U-shape or parabolic trend, which implies the model has the lack of fit issue. We might consider adding the second order terms of predictors or use partial residual plots to get more information on the relationship between y and the predictors.


Question 6

Use the following code to plot the partial residual plots. Comment on the partial residual plots. What do the plots reveal the information between medv and the independent variables?

```
library(car)
crPlots(fit)
```



Component + Residual Plots



Answer: From the partial residual plots above, we could find the medv vs rm, and medv vs lstat has significant curvature.

We should consider adding the quadratic terms of rm and lstat.

Question 7

Fit the model with 10 variables in Question 5 and add second-order term of lstat, second-order term of rm, and the interaction between rad & lstat, rm & rad, crim & chas, chas & nox.

```
fit.c = lm(medv~crim + chas + nox + rm + dis + rad + tax + ptratio + black + lstat
           + I(lstat^2) + I(rm^2) + rad:lstat + rm:rad + crim:chas + chas:nox,
           data = Boston)
summary(fit.c)
```

```
##
## Call:
## lm(formula = medv ~ crim + chas + nox + rm + dis + rad + tax +
##      ptratio + black + lstat + I(lstat^2) + I(rm^2) + rad:lstat +
##      rm:rad + crim:chas + chas:nox, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.610  -2.026  -0.223   1.604  24.638
```




```
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  42.046646  11.485042   3.661 0.000279 ***
## crim        -0.111171   0.024327  -4.570 6.19e-06 ***
## chas         13.473038   3.087684   4.363 1.56e-05 ***
## nox        -12.363598   2.764480  -4.472 9.63e-06 ***
## rm          -4.427683   3.056792  -1.448 0.148126
## dis         -0.820218   0.120105  -6.829 2.54e-11 ***
## rad          2.802634   0.264222  10.607 < 2e-16 ***
## tax         -0.008726   0.002397  -3.640 0.000301 ***
## ptratio     -0.745128   0.090980  -8.190 2.29e-15 ***
## black        0.003524   0.001948   1.809 0.071130 .
## lstat       -0.881908   0.109460  -8.057 6.04e-15 ***
## I(lstat^2)    0.024918   0.003128   7.967 1.15e-14 ***
## I(rm^2)       0.973805   0.216680   4.494 8.73e-06 ***
## rad:lstat    -0.035154   0.003714  -9.466 < 2e-16 ***
## rm:rad       -0.338581   0.037780  -8.962 < 2e-16 ***
## crim:chas     2.141012   0.322928   6.630 8.90e-11 ***
## chas:nox    -26.028896   5.684180  -4.579 5.93e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.41 on 489 degrees of freedom
## Multiple R-squared:  0.8669, Adjusted R-squared:  0.8625
## F-statistic: 199 on 16 and 489 DF, p-value: < 2.2e-16
```

```
source("anova_alt.R")
anova_alt(fit.c)
```

```
## Analysis of Variance Table
##
##           Df      SS      MS      F      P
## Source   16 37030 2314.36 199.01 3.2656e-202
## Error   489   5687   11.63
## Total   505 42716   84.59
```

```
par(mfrow=c(2,2))
plot(fit.c)
```



```
par(mfrow=c(1,1))
```

Comment on the model fitted.

Answer: From the output above, we could see that overall F-test yield significant conclusion. What's more, the $R_a^2 = 86.25\%$ and $SE = 3.41$. The overall model is statistically useful for predicting medv.

We then examine the residual plots.

- From the residual plot, there is no significant patterns residual plots in the upper left corner.
- The Normal QQ plot generally looks OK.
- The plot in the lower left corner doesn't show significant increasing or decreasing trend, which suggests homogeneous variance. By using standardized residual method, there are outliers.
- Observation 365 is influential.

Question 8

For the model fitted in Question 7, is the normality assumption reasonably satisfied?

Answer: Yes, based on the Normal QQ plot in Question 7, the normality assumption is reasonably satisfied.

Question 9

Use Studentized Deleted Residual method to identify whether there are any outliers.

```
studresid = rstudent(fit.c)
studresid_sorted = sort(studresid)
head(round(studresid_sorted, 3), 10)
```

```
##      365      357      366      89      506      427      343      309      11      398
## -5.863 -3.372 -2.775 -2.530 -2.437 -2.129 -2.104 -1.887 -1.814 -1.809
```

```
tail(round(studresid_sorted, 3), 10)
```

```
##      370      153      392      187      162      182      373      410      369      372
##  2.261  2.365  2.481  2.567  2.802  3.365  3.583  4.022  4.281  7.780
```

Answer: From the R output above, we could find there are several outliers: Observation #365, #357, #182, #373, #410, #369, #372, which all have the studentized deleted residuals with absolute value greater than 3.

Our model doesn't contain all the significant interactions. This might be the reason that there are several outliers.

Question 10

Use Cook's Distance method to identify whether there are any influential observations.

```
cookdist = cooks.distance(fit.c)
cooks_sorted = sort(round(cookdist, 3))

tail(round(cooks_sorted, 3), 10)
```

```
##      413      155      370      372      153      366      369      373      357      365
##  0.032  0.034  0.053  0.107  0.119  0.153  0.158  0.307  0.320  1.014
```

Answer: The model fitted in Question 7 has 16 independent variables, thus $k = 16$. There are 506 observations. Thus, $n = 506$.

Thus $\nu_1 = 16 + 1$, $\nu_2 = n - (k + 1) = 506 - (16 + 1) = 489$. Therefore,

$$F_{0.5, 17, 489} = 0.96$$

Therefore, observation 365 has cook's distance = 1.014 > 0.96. We have observation 365 is an influential observation. This conclusion is consistent with the conclusion in Question 7.